

From Testing to Proof using Symbolic Execution

Aaron Tomb
Galois, Inc.

StrangeLoop
September 28, 2017

- Installation (~15m)
- Basic overview (~15m)
- Exercises (~20m)
- More flexible verification (~15m)
- Exercises (~20m)
- Composition and more (~15m)
- Exercises (~20m)

- SAW 2017-09-06 from <https://saw.galois.com/builds/nightly/>
- Z3 4.5.0 from <https://github.com/Z3Prover/z3/releases>
- SAW builds available for:
 - ▶ CentOS 6 (32-bit and 64-bit) (anything similar to older RedHat)
 - ▶ CentOS 7 (64-bit) (anything similar to newer RedHat)
 - ▶ macOS (64-bit)
 - ▶ Ubuntu 14.04 (64-bit) (anything similar to recent-ish Debian)
 - ▶ Windows (64-bit)
- VirtualBox VM image and local(ish) tarballs:
 - ▶ <http://10.129.176.174:8000/>
 - ▶ SAW Workshop (Debian).vdi
 - ▶ saw/* (files for various platforms)
 - ▶ z3/* (files for various platforms)
 - ▶ Login: root/saw-workshop, saw/saw-workshop

- A tool to construct **models** of program behavior
 - ▶ Works with C (LLVM), Java (JVM), and others in progress
 - ▶ Also supports specifications written in Cryptol
- Models can then be **proved** to have certain properties
 - ▶ Equivalence with specifications
 - ▶ Guarantees to return certain values
- Proofs generally done using **automated** reasoning tools
 - ▶ So similar level of effort to testing
 - ▶ Uses a technique called symbolic execution, plus SAT/SMT

- Rather than testing individual cases, state general properties
- Then can test those properties on specific values
 - ▶ Manually selected
 - ▶ Randomly generated
- For example, this function should always return a non-zero value:

```
int add_commutes(uint32_t x, uint32_t y) {  
    return x + y == y + x;  
}
```

- The QuickCheck approach is a common implementation of this paradigm

- Say we're using the XOR-based trick for swapping values:

```
void swap_xor(uint32_t *x, uint32_t *y) {  
    *x = *x ^ *y;  
    *y = *x ^ *y;  
    *x = *x ^ *y;  
}
```

- Focus on values, since that's where the tricky parts are
 - ▶ Pointers used just so it can be a separate function

```
void swap_direct(uint32_t *x, uint32_t *y) {  
    uint32_t tmp;  
    tmp = *y;  
    *y = *x;  
    *x = tmp;  
}
```

```
int swap_correct(uint32_t x, uint32_t y) {  
    uint32_t x1 = x, x2 = x, y1 = y, y2 = y;  
    swap_xor(&x1, &y1);  
    swap_direct(&x2, &y2);  
    return (x1 == x2 && y1 == y2);  
}
```

```
int main() {  
    assert(swap_correct(0, 0));  
    assert(swap_correct(0, 1));  
    assert(swap_correct(1, 0));  
    assert(swap_correct(32, 76));  
    assert(swap_correct(0, 0xFFFFFFFF));  
    assert(swap_correct(0xFFFFFFFF, 0xFFFFFFFF));  
    return 0;  
}
```

- Advantages
 - ▶ Ensures that you will always test important values
 - ▶ Carefully chosen tests can cover many important cases quickly
- Disadvantages
 - ▶ May miss classes of inputs that you didn't think of
 - ▶ Non-deterministics: different runs may have different results


```
int main() {  
    for(int idx = 0; i < 100; i++) {  
        uint32_t x = rand();  
        uint32_t y = rand();  
        assert(swap_correct(x, y));  
    }  
    return 0;  
}
```

- Advantages

- ▶ Better theoretical coverage of input space
- ▶ Number of tests limited only by available processing power

- Disadvantages

- ▶ May miss important classes of inputs that are easy to identify by hand

Translating Programs to Pure Functions

- $\lambda x. x + 1$ is a function
 - ▶ takes an argument x , and returns $x + 1$
- `swap_direct`: $\lambda(x, y). (y, x)$
- `swap_xor`: $\lambda(x, y). (x \oplus y \oplus x \oplus y \oplus y, x \oplus y \oplus y)$
 - ▶ but $x \oplus x \equiv 0$ and $x \oplus 0 \equiv x$
- Translation achieved in SAW using a technique called *symbolic execution*
 - ▶ Think: an interpreter with expressions in place of values
 - ▶ Every variable's value at the end is an expression representing *all possible values* it might take

- Automated provers for mathematical theorems
 - ▶ Such as: $\forall x, y. (x \oplus y \oplus x \oplus y \oplus y, x \oplus y \oplus y) \equiv (y, x)$
- SAT = Boolean SATisfiability
- SMT = Satisfiability Modulo Theories
- Almost magic for what they can do. SAT can encode:
 - ▶ Fixed-size bit vectors (even multiplication, but slowly)
 - ▶ Bit manipulation operations (and, or, xor, shifts)
 - ▶ Arrays of fixed sizes
 - ▶ Conditionals
- SMT adds things like:
 - ▶ Linear arithmetic on integers (addition, subtraction, multiplication by constants)
 - ▶ Arrays of arbitrary size
 - ▶ Uninterpreted functions

- Advantages
 - ▶ Ensures that you will test **all possible** input values
 - ▶ Sometimes faster than testing
- Disadvantages
 - ▶ Applicable to a smaller class of programs than testing
 - ▶ Sometimes much slower than testing

```
// Load the bitcode file generated by Clang
swapmod <- llvm_load_module "swap.bc";

// Extract a formal model of `swap_correct`
harness <- llvm_extract swapmod "swap_correct" llvm_pure;

// Use ABC prover to show it always returns non-zero
prove_print abc {{ \x y -> harness x y != 0 }};

(ln swap_harness.saw)
```

```
uint32_t ffs_ref(uint32_t word) {  
    if(!word) return 0;  
    for(int c = 0, i = 0; c < 32; c++)  
        if(((1 << i++) & word) != 0)  
            return i;  
    return 0;  
}
```

```
uint32_t ffs_imp(uint32_t i) {  
    char n = 1;  
    if (!(i & 0xffff)) { n += 16; i >>= 16; }  
    if (!(i & 0x00ff)) { n += 8; i >>= 8; }  
    if (!(i & 0x000f)) { n += 4; i >>= 4; }  
    if (!(i & 0x0003)) { n += 2; i >>= 2; }  
    return (i) ? (n+((i+1) & 0x01)) : 0;  
}
```

```
int ffs_imp_correct(uint32_t x) {  
    return ffs_imp(x) == ffs_ref(x);  
}  
  
int main() {  
    assert(ffs_imp_correct(0x00000000));  
    assert(ffs_imp_correct(0x00000001));  
    assert(ffs_imp_correct(0x80000000));  
    assert(ffs_imp_correct(0x80000001));  
    assert(ffs_imp_correct(0xF0000000));  
    assert(ffs_imp_correct(0x0000000F));  
    assert(ffs_imp_correct(0xFFFFFFFF));  
    return 0;  
}
```

- Same pros and cons as for the swap example

```
int main() {  
    for(int idx = 0; i < 100; i++) {  
        uint32_t x = rand();  
        assert(ffs_imp_correct(x));  
    }  
    return 0;  
}
```



```
m <- llvm_load_module "ffs.bc";

correct <- llvm_extract m "ffs_imp_correct" llvm_pure;

print "Proving ffs_imp_correct always returns true...";
prove_print abc {{ \x -> correct x == 1 }};

(ln ffs_harness.saw)
```

```
m <- llvm_load_module "ffs.bc";

ref <- llvm_extract m "ffs_ref" llvm_pure;
imp <- llvm_extract m "ffs_imp" llvm_pure;

// Following equivalent to \x -> ref x == imp x
prove_print abc {{ ref === imp }};

(ln ffs_eq.saw)
```

0. Run the equivalence proofs in `ffs_harness.saw` and `ffs_eq.saw`
1. Port the FFS code to use `uint64_t`
 - Translate both reference and implementation
 - Which one is wrong?
2. Try to break the FFS code, in obvious and subtle ways
 - Can you make it do the wrong thing and not be caught?
3. Try to discover the “haystack” bug in `ffs_bug`
 - Use random testing (`ffs_bug_fail.saw`)
 - ▶ Increase the number of tests and see how long it takes
 - ▶ Try a similar case with `uint64_t`
 - Use `ffs_bug.saw` to find it with a SAT solver

Pointers: Verifying XOR Swap Without Wrapper

```
m <- llvm_load_module "xor-swap.bc";
// void swap_xor(uint32_t *x, uint32_t *y);
let swap_spec = do {
  x <- crucible_fresh_var "x" (llvm_int 32);
  y <- crucible_fresh_var "y" (llvm_int 32);
  xp <- crucible_alloc (llvm_int 32);
  yp <- crucible_alloc (llvm_int 32);
  crucible_points_to xp (crucible_term x);
  crucible_points_to yp (crucible_term y);
  crucible_execute_func [xp, yp];
  crucible_points_to xp (crucible_term y);
  crucible_points_to yp (crucible_term x);
};
crucible_llvm_verify m "swap_xor" [] true swap_spec abc;

(In swap.saw)
```

Simplifying the XOR Swap specification

```
m <- llvm_load_module "xor-swap.bc";

let ptr_to_fresh nm ty = do {
  x <- crucible_fresh_var nm ty;
  p <- crucible_alloc ty;
  crucible_points_to p (crucible_term x);
  return (x, p);
};

let swap_spec = do {
  (x, xp) <- ptr_to_fresh "x" (llvm_int 32);
  (y, yp) <- ptr_to_fresh "y" (llvm_int 32);
  crucible_execute_func [xp, yp];
  crucible_points_to xp (crucible_term y);
  crucible_points_to yp (crucible_term x);
};
```

1. Try to break the XOR-based swapping in some way and run the proof
 - ▶ Use `swap.saw` or `swap_harness.saw`
2. Write a buggy version and use SAW to find inputs for which it's correct
 - These would be bad test cases!
3. Write a script to prove the FFS test harness using `crucible_llvm_verify`
 - You'll need `crucible_return {{ 1 : [32] }}` and `crucible_term`
 - You won't need `crucible_alloc` or `crucible_points_to`

- Verifications in SAW consist of three phases
 - ▶ Initialize a starting state
 - ▶ Run the target code in that state
 - ▶ Check that the final state is correct
- Commands like `llvm_extract` just simplify a common case
- When running the target code, we can sometimes use previously-proven facts about code it calls

Composition: Verifying Salsa20 (C code)

```
uint32_t rotl(uint32_t value, int shift) {  
    return (value << shift) | (value >> (32 - shift));  
}
```

```
void s20_quarterround(uint32_t *y0, uint32_t *y1,  
                      uint32_t *y2, uint32_t *y3) {  
    *y1 = *y1 ^ rotl(*y0 + *y3, 7);  
    // ... and three more  
}
```

```
void s20_rowround(uint32_t y[static 16]) {  
    s20_quarterround(&y[0], &y[1], &y[2], &y[3]);  
    // ... and three more  
}
```


Composition: Verifying Salsa20 (SAW code)

```
let quarterround_setup : CrucibleSetup () = do {  
  (p0, y0) <- ptr_to_fresh "y0" i32;  
  // ... and three more  
  crucible_execute_func [p0, p1, p2, p3];  
  let zs = [{ quarterround [y0,y1,y2,y3] }];  
  crucible_points_to p0 (crucible_term [{ zs@0 }]);  
  // ... and three more  
};  
  
let rowround_setup = do {  
  (y, p) <- ptr_to_fresh "y" (llvm_array 16 i32);  
  crucible_execute_func [p];  
  crucible_points_to p (crucible_term [{ rowround y }]);  
};
```

- With the current version of SAW, programs must be **finite**
 - ▶ SAT-based proofs need to know how many bits are involved
 - ▶ Inputs need to have fixed sizes
 - ▶ All pointers must point to data of known size
 - ▶ All loops need to execute a bounded number of types
- But Salsa20 can operate on any input size
 - ▶ So we prove it correct separately for several possible sizes
 - ▶ Our original version had a bug because of this!
- Future versions are likely to relax these restrictions

1. Run the monolithic and compositional proofs
 - `salsa.saw` and `salsa-compositional.saw`
2. Compare the timing of the two
 - When checking multiple sizes, how does it compare?
 - How many sizes before it becomes better?
3. Try to break the code and see what happens
 - First try a leaf function
 - Then try the top-level function
4. Can you break it so that one size succeeds but another fails?

Sidebar: Fuzzing for Property Based Tests

- Klee is another LLVM symbolic execution system
- Doesn't aim for complete coverage, but very powerful
 - ▶ Better at finding bugs than SAW
 - ▶ Not generally usable for verification
- Associated fuzz testing system, libfuzzer
 - ▶ Includes a main function that calls fuzzing harness

```
int LLVMFuzzerTestOneInput(const uint8_t *Data,
                           size_t Size) {
    DoSomethingInterestingWithMyAPI(Data, Size);
    return 0;
}
```

```
let fuzzer_spec n = do {  
  let ty = llvm_array n i8;  
  (pdata, data) <- ptr_to_fresh "data" ty;  
  crucible_execute_func  
    [ pdata  
      , crucible_term {{ `n : [64] }}  
    ];  
  crucible_return (crucible_term {{ 0 : [32] }});  
};  
  
m <- llvm_load_module "fuzztarget.bc";  
for [1, 10, 20, 100] (\sz ->  
  crucible_llvm_verify m "LLVMFuzzerTestOneInput"  
    [] true (fuzzer_spec n) abc);
```

- Support for various languages
 - ▶ Others that compile to LLVM (simple C++ and Rust have been tested, others YMMV)
 - ▶ Languages that compile to JVM (only Java known to work well, but others might)
 - ▶ Coming soon: Rust, Go, some degree of machine code
- Some interactive proof tactics
 - ▶ Mostly rewriting with user-defined rules
 - ▶ Coming soon: bindings to external interactive provers, including Lean and Coq

- Resources
 - ▶ SAW web site: <https://saw.galois.com>
 - ▶ Cryptol web site: <https://cryptol.net>
 - ▶ SAW documentation
 - ▶ Tutorial: <https://saw.galois.com/tutorial.html>
 - ▶ Manual: <https://saw.galois.com/manual.html>
 - ▶ Cryptol documentation:
<https://cryptol.net/documentation.html>
- I'll be around all day, and happy to talk more.
- And if this sort of thing interests you, Galois is hiring!